

SGN-41007 Pattern Recognition and Machine Learning

Exercise Set 7: February 18–February 22, 2018

Exercises consist of both pen&paper and computer assignments. Pen&paper questions are solved at home before exercises, while computer assignments are solved during exercise hours. The computer assignments are marked by `python` and Pen&paper questions by `pen&paper`

1. `pen&paper` *Error rate confidence limits.*

We train a classifier with a set of training examples, and test the accuracy of the resulting model with a set of $N = 100$ test samples. The classifier misclassifies $K = 5$ of those.

- a) Find the 90% confidence interval of the result. Hint: The classification accuracy can be modeled using binomial distribution, whose confidence intervals are discussed here:

https://en.wikipedia.org/wiki/Binomial_distribution#Confidence_intervals

- b) Another classifier misclassifies only 3 test samples. Is it better than the first one with statistical significance at 90% confidence level?

2. `pen&paper` *Design a regularized LDA classifier.*

Let's revisit the LDA design of Exercise set 4, but add a regularization term. The non-regularized LDA solution is given by as

$$\mathbf{w} = (\Sigma_0 + \Sigma_1)^{-1} (\mu_1 - \mu_0)$$

The regularized solution with regularization parameter $\lambda > 0$ is defined as

$$\mathbf{w} = (\Sigma_0 + \Sigma_1 + \lambda \mathbf{I})^{-1} (\mu_1 - \mu_0)$$

However, as the scale of \mathbf{w} is not important—only the direction—let us use an alternative definition instead:

$$\mathbf{w} = \lambda (\Sigma_0 + \Sigma_1 + \lambda \mathbf{I})^{-1} (\mu_1 - \mu_0).$$

This definition avoids the convergence of \mathbf{w} towards zero as $\lambda \rightarrow \infty$.

- a) Compute the regularized LDA weight vector¹ for $\lambda = 100$ and

$$\begin{aligned} \mu_0 &= \begin{pmatrix} 1 \\ 1 \end{pmatrix} & \mu_1 &= \begin{pmatrix} 0 \\ 0 \end{pmatrix} \\ \Sigma_0 &= \begin{pmatrix} 3 & -2 \\ -2 & 2 \end{pmatrix} & \Sigma_1 &= \begin{pmatrix} 3 & -2 \\ -2 & 2 \end{pmatrix} \end{aligned}$$

- b) Where does \mathbf{w} converge as $\lambda \rightarrow \infty$?

¹Remember the inversion rule for 2×2 matrices:

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix}^{-1} = \frac{1}{ad - bc} \begin{pmatrix} d & -b \\ -c & a \end{pmatrix}$$

3. **python** Let us use a pretrained VGG16 model for last week's GTSRB experiment. Instead of using the custom ConvNet of last week, initialize a VGG16 net as described on Slide 15 of the slide set at <http://www.cs.tut.fi/courses/SGN-41007/slides/Lecture6b.pdf>. Add dense layers after the convolutional pipeline such that `model.summary()` reports the following (top of listing omitted):

<code>block5_conv3</code> (Conv2D)	(None, 4, 4, 512)	2359808
<code>block5_pool</code> (MaxPooling2D)	(None, 2, 2, 512)	0
<code>flatten_1</code> (Flatten)	(None, 2048)	0
<code>dense_3</code> (Dense)	(None, 100)	204900
<code>dense_4</code> (Dense)	(None, 2)	202
=====		
<code>Total params:</code> 14,919,790		
<code>Trainable params:</code> 14,919,790		
<code>Non-trainable params:</code> 0		

Compile and run the net. Note that you will need a GPU (e.g., TC303 machines) for training this net.

4. **python** Apply the recursive feature elimination approach (`sklearn.feature_selection.RFECV`) with logistic regression classifier for the arcene dataset. The data can be downloaded in *.mat format from:

<http://www.cs.tut.fi/courses/SGN-41007/exercises/arcene.zip>

Use `scipy.io.loadmat` to open the file. Note that you have to ravel `y_train` and `y_test` so that `sklearn` will accept them.

- Instantiate an RFECV selector (call it `rfe` from now on). To speed up computation, set `step = 50` in the constructor. Also set `verbose = 1` to see the progress.
 - Fit the RFECV to `X_train` and `y_train`.
 - Count the number of selected features from `rfe.support_`.
 - Plot the errors for different number of features:
`plt.plot(range(0,10001,50), rfe.grid_scores_)`
 - Compute the accuracy on `X_test` and `y_test`. You can use `rfe` as any other classifier.
5. **python** Apply L_1 penalized Logistic Regression for feature selection with the arcene dataset. Find a good value for parameter C by 10-fold cross-validating the accuracy. Study the sparseness of the solution: how many features were selected?
- Instantiate a LogisticRegression classifier. Set `penalty = 'l1'` in the constructor.

- b) Cross validate the accuracy of a range of `C` values (see earlier exercises).
- c) Fit the `LogisticRegression` to `X_train` and `y_train`.
- d) Count the number of selected features from `clf.coef_`, where `clf` is your logistic regression classifier.
- e) Compute the accuracy on `X_test` and `y_test`.